

Тема: Методы построения моделей машинного обучения

Загрузка необходимых библиотек:

```
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn import preprocessing
from sklearn import svm
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_squared_log_error, median_absolute_error, r2_score
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.ensemble import RandomForestClassifier

from sklearn.ensemble import AdaBoostClassifier
from sklearn.impute import SimpleImputer
```

```
data = pd.read_csv('./HRDataset_v14.csv', sep=",")
TARGET_COL_NAME = 'RecruitmentSource'
TARGET_IS_NUMERIC = data[TARGET_COL_NAME].dtype != 'O'
TARGET_IS_NUMERIC
```

False

data

[illegible]

306	Employee Name	EmpID	MarriedID	MaritalStatusID	GenderID	EmpStatusID	DeptID	PerfScoreID	FromDiversityJobFair
307	Ybarra, Catherine	10301	0	0	0	5	5	1	
308	Zamora, Jennifer	10010	0	0	0	1	3	4	
309	Zhou, Julia	10043	0	0	0	1	3	3	
310	Zima, Colleen	10271	0	4	0	1	5	3	

311 rows x 36 columns



In [4]:

```
data.shape
```

Out[4]:

(311, 36)

In [5]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 311 entries, 0 to 310
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Employee_Name                        311 non-null    object
1   EmpID                               311 non-null    int64
2   MarriedID                           311 non-null    int64
3   MaritalStatusID                     311 non-null    int64
4   GenderID                            311 non-null    int64
5   EmpStatusID                         311 non-null    int64
6   DeptID                              311 non-null    int64
7   PerfScoreID                         311 non-null    int64
8   FromDiversityJobFairID              311 non-null    int64
9   Salary                              311 non-null    int64
10  Termd                               311 non-null    int64
11  PositionID                           311 non-null    int64
12  Position                             311 non-null    object
13  State                                311 non-null    object
14  Zip                                  311 non-null    int64
15  DOB                                  311 non-null    object
16  Sex                                  311 non-null    object
17  MaritalDesc                          311 non-null    object
18  CitizenDesc                          311 non-null    object
19  HispanicLatino                       311 non-null    object
20  RaceDesc                             311 non-null    object
21  DateofHire                           311 non-null    object
22  DateofTermination                    104 non-null    object
23  TermReason                           311 non-null    object
24  EmploymentStatus                     311 non-null    object
25  Department                           311 non-null    object
26  ManagerName                          311 non-null    object
27  ManagerID                            303 non-null    float64
28  RecruitmentSource                    311 non-null    object
29  PerformanceScore                     311 non-null    object
30  EngagementSurvey                     311 non-null    float64
31  EmpSatisfaction                      311 non-null    int64
32  SpecialProjectsCount                 311 non-null    int64
33  LastPerformanceReview_Date           311 non-null    object
34  DaysLateLast30                       311 non-null    int64
35  Absences                             311 non-null    int64
dtypes: float64(2), int64(16), object(18)
memory usage: 87.6+ KB
```

In [6]:

```
# проверим есть ли пропущенные значения
data.isnull().sum()
```

Out[6]:


```
Employee_Name          0
EmpID                  0
MarriedID              0
MaritalStatusID        0
GenderID               0
EmpStatusID            0
DeptID                 0
PerfScoreID            0
FromDiversityJobFairID 0
Salary                 0
Termd                  0
PositionID             0
Position               0
State                  0
Zip                    0
DOB                    0
Sex                    0
MaritalDesc            0
CitizenDesc            0
HispanicLatino         0
RaceDesc               0
DateofHire              0
DateofTermination      207
TermReason             0
EmploymentStatus       0
Department             0
ManagerName            0
ManagerID              8
RecruitmentSource      0
PerformanceScore       0
EngagementSurvey       0
EmpSatisfaction         0
SpecialProjectsCount    0
LastPerformanceReview_Date 0
DaysLateLast30         0
Absences               0
dtype: int64
```

Удалим колонки, которые не влияют на целевой признак

Построим **heatmap** для лучшего визуального представления всех корреляций

In [7]:

```
cmap = sns.cm.rocket_r
ax = sns.heatmap(data.corr(), cmap=cmap)
```



In [8]:

```
data = data.drop(columns=['Employee_Name', 'EmpID', 'DateofTermination', 'ManagerID'])
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 311 entries, 0 to 310
Data columns (total 32 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   MarriedID                                311 non-null    int64
1   MaritalStatusID                         311 non-null    int64
2   GenderID                                311 non-null    int64
3   EmpStatusID                             311 non-null    int64
4   DeptID                                  311 non-null    int64
5   PerfScoreID                             311 non-null    int64
6   FromDiversityJobFairID                  311 non-null    int64
7   Salary                                  311 non-null    int64
8   Termd                                   311 non-null    int64
9   PositionID                              311 non-null    int64
10  Position                                 311 non-null    object
11  State                                   311 non-null    object
12  Zip                                     311 non-null    int64
13  DOB                                    311 non-null    object
14  Sex                                    311 non-null    object
15  MaritalDesc                             311 non-null    object
16  CitizenDesc                             311 non-null    object
17  HispanicLatino                          311 non-null    object
18  RaceDesc                                311 non-null    object
19  DateofHire                              311 non-null    object
20  TermReason                              311 non-null    object
21  EmploymentStatus                        311 non-null    object
22  Department                              311 non-null    object
23  ManagerName                             311 non-null    object
24  RecruitmentSource                       311 non-null    object
25  PerformanceScore                        311 non-null    object
26  EngagementSurvey                        311 non-null    float64
27  EmpSatisfaction                         311 non-null    int64
28  SpecialProjectsCount                    311 non-null    int64
29  LastPerformanceReview_Date              311 non-null    object
30  DaysLateLast30                          311 non-null    int64
31  Absences                                311 non-null    int64
dtypes: float64(1), int64(15), object(16)
memory usage: 77.9+ KB
```

Обработка пропусков

In [9]:

```
# Импутация наиболее частыми значениями
imp = SimpleImputer(missing_values=np.nan, strategy='most_frequent')

imputed = {}

for col in data:
    contains_nan = data[col].isnull().sum() != 0
    if contains_nan:
        data_imp = data[[col]]
        data_imp = imp.fit_transform(data_imp)
        imputed[col] = data_imp

for col_name in imputed:
    df = pd.DataFrame({col_name:imputed[col_name].T[0]})
    data[col_name] = df.copy()

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 311 entries, 0 to 310
Data columns (total 32 columns):
#      Column                                     Non-Null Count  Dtype
---  -
0      MarriedID                                     311 non-null    int64
1      MaritalStatusID                             311 non-null    int64
2      GenderID                                     311 non-null    int64
3      EmpStatusID                                 311 non-null    int64
4      DeptID                                       311 non-null    int64
5      PerfScoreID                                311 non-null    int64
6      FromDiversityJobFairID                      311 non-null    int64
7      Salary                                       311 non-null    int64
8      Termd                                       311 non-null    int64
9      PositionID                                 311 non-null    int64
10     Position                                    311 non-null    object
11     State                                       311 non-null    object
12     Zip                                         311 non-null    int64
13     DOB                                         311 non-null    object
14     Sex                                         311 non-null    object
15     MaritalDesc                               311 non-null    object
16     CitizenDesc                               311 non-null    object
17     HispanicLatino                           311 non-null    object
18     RaceDesc                                  311 non-null    object
19     DateofHire                                311 non-null    object
20     TermReason                                311 non-null    object
21     EmploymentStatus                         311 non-null    object
22     Department                                311 non-null    object
23     ManagerName                              311 non-null    object
24     RecruitmentSource                        311 non-null    object
25     PerformanceScore                         311 non-null    object
26     EngagementSurvey                        311 non-null    float64
27     EmpSatisfaction                          311 non-null    int64
28     SpecialProjectsCount                    311 non-null    int64
29     LastPerformanceReview_Date              311 non-null    object
30     DaysLateLast30                          311 non-null    int64
31     Absences                                311 non-null    int64
dtypes: float64(1), int64(15), object(16)
memory usage: 77.9+ KB
```

Кодирование строковых признаков (LabelEncoding)

In [10]:

```
not_number_cols = data.select_dtypes(include=['object'])
number_cols = data.select_dtypes(exclude=['object'])
```

In [11]:

```
le = preprocessing.LabelEncoder()

for col_name in not_number_cols:
    data[col_name] = le.fit_transform(data[col_name])

data
```

Out[11]:

	MarriedID	MaritalStatusID	GenderID	EmpStatusID	DeptID	PerfScoreID	FromDiversityJobFairID	Salary	Termd	Positic
0	0	0	1	1	5	4	0	62506	0	
1	1	1	1	5	3	3	0	104437	1	
2	1	1	0	5	5	3	0	64955	1	
3	1	1	0	1	5	3	0	64991	0	
4	0	2	0	5	5	3	0	50825	1	
...	
306	0	0	1	1	5	3	0	65893	0	

	MarriedID	MaritalStatusID	GenderID	EmpStatusID	DeptID	PerfScoreID	FromDiversityJobFairID	Salary	Termd	Positi
307	0	0	0	5	5	1	0	48513	1	
308	0	0	0	1	3	4	0	220450	0	
309	0	0	0	1	3	3	0	89292	0	
310	0	4	0	1	5	3	0	45046	0	

311 rows x 32 columns



Масштабируем числовые данные

In [12]:

```
scaler = preprocessing.MinMaxScaler()

number_fields_source = number_cols.loc[:, number_cols.columns!=TARGET_COL_NAME] if TARGE
T_IS_NUMERIC else number_cols

for col_name in number_fields_source:
    data[col_name] = scaler.fit_transform(data[[col_name]])

data
```

Out[12]:

	MarriedID	MaritalStatusID	GenderID	EmpStatusID	DeptID	PerfScoreID	FromDiversityJobFairID	Salary	Termd	Positi
0	0.0	0.00	1.0	0.0	0.8	1.000000	0.0	0.085190	0.0	0.6
1	1.0	0.25	1.0	1.0	0.4	0.666667	0.0	0.289777	1.0	0.8
2	1.0	0.25	0.0	1.0	0.8	0.666667	0.0	0.097139	1.0	0.6
3	1.0	0.25	0.0	0.0	0.8	0.666667	0.0	0.097315	0.0	0.6
4	0.0	0.50	0.0	1.0	0.8	0.666667	0.0	0.028197	1.0	0.6
...
306	0.0	0.00	1.0	0.0	0.8	0.666667	0.0	0.101716	0.0	0.6
307	0.0	0.00	0.0	1.0	0.8	0.000000	0.0	0.016916	1.0	0.6
308	0.0	0.00	0.0	0.0	0.4	1.000000	0.0	0.855821	0.0	0.1
309	0.0	0.00	0.0	0.0	0.4	0.666667	0.0	0.215883	0.0	0.2
310	0.0	1.00	0.0	0.0	0.8	0.666667	0.0	0.000000	0.0	0.6

311 rows x 32 columns



Делим выборку на обучающую и тестовую

In [13]:

```
target = data[TARGET_COL_NAME]
data_X_train, data_X_test, data_y_train, data_y_test = train_test_split(
    data, target, test_size=0.2, random_state=1)
```

In [14]:

```
data_X_train.shape, data_y_train.shape
```

Out[14]:

((248, 32), (248,))

In [15]:

```
data_X_test.shape, data_y_test.shape
```

```
Out[15]:  
  
((63, 32), (63,))
```

```
In [16]:
```

```
np.unique(target)
```

```
Out[16]:  
  
array([0, 1, 2, 3, 4, 5, 6, 7, 8])
```

Логистическая регрессия

```
In [17]:
```

```
svr_1 = LogisticRegression(solver='lbfgs', max_iter=1000)  
svr_1.fit(data_X_train, data_y_train)
```

C:\Users\pstri\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
Out[17]:  
  
LogisticRegression(max_iter=1000)
```

```
In [18]:
```

```
data_y_pred_1 = svr_1.predict(data_X_test)  
accuracy_score(data_y_test, data_y_pred_1)
```

```
Out[18]:  
  
0.7619047619047619
```

```
In [19]:
```

```
f1_score(data_y_test, data_y_pred_1, average='micro')
```

```
Out[19]:  
  
0.7619047619047619
```

```
In [20]:
```

```
f1_score(data_y_test, data_y_pred_1, average='macro')
```

```
Out[20]:  
  
0.7208312792201521
```

```
In [21]:
```

```
f1_score(data_y_test, data_y_pred_1, average='weighted')
```

```
Out[21]:  
  
0.7563589699202566
```

```
In [22]:
```

```
svr_2 = LogisticRegression(solver='lbfgs', max_iter=10000)  
svr_2.fit(data_X_train, data_y_train)
```

```
C:\Users\pstri\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
```

```
Out[22]:
```

```
LogisticRegression(max_iter=10000)
```

```
In [23]:
```

```
data_y_pred_2 = svr_2.predict(data_X_test)
accuracy_score(data_y_test, data_y_pred_2)
```

```
Out[23]:
```

```
0.8095238095238095
```

```
In [24]:
```

```
f1_score(data_y_test, data_y_pred_2, average='micro')
```

```
Out[24]:
```

```
0.8095238095238095
```

```
In [25]:
```

```
f1_score(data_y_test, data_y_pred_2, average='macro')
```

```
Out[25]:
```

```
0.5976659982174688
```

```
In [26]:
```

```
f1_score(data_y_test, data_y_pred_2, average='weighted')
```

```
Out[26]:
```

```
0.7990160710748947
```

Случайный лес

```
In [27]:
```

```
RT = RandomForestClassifier(n_estimators=15, random_state=123)
RT.fit(data_X_train, data_y_train)
```

```
Out[27]:
```

```
RandomForestClassifier(n_estimators=15, random_state=123)
```

```
In [28]:
```

```
accuracy_score(data_y_test, RT.predict(data_X_test))
```

```
Out[28]:
```

```
0.7777777777777778
```

```
In [29]:
```

```
f1_score(data_y_test, data_y_pred_1, average='micro')
```

```
Out[29]:
```

```
0.7619047619047619
```


In [30]:

```
f1_score(data_y_test, data_y_pred_1, average='macro')
```

Out[30]:

0.7208312792201521

In [31]:

```
f1_score(data_y_test, data_y_pred_1, average='weighted')
```

Out[31]:

0.7563589699202566

In [32]:

```
RT = RandomForestClassifier(n_estimators=30, random_state=123)
RT.fit(data_X_train, data_y_train)
```

Out[32]:

RandomForestClassifier(n_estimators=30, random_state=123)

In [33]:

```
accuracy_score(data_y_test, RT.predict(data_X_test))
```

Out[33]:

0.873015873015873

In [34]:

```
f1_score(data_y_test, data_y_pred_1, average='micro')
```

Out[34]:

0.7619047619047619

In [35]:

```
f1_score(data_y_test, data_y_pred_1, average='macro')
```

Out[35]:

0.7208312792201521

In [36]:

```
f1_score(data_y_test, data_y_pred_1, average='weighted')
```

Out[36]:

0.7563589699202566

Выводы

При использовании логистической регрессии наилучшую точность **(0.809)** показала модель с параметром `max_iter=10000`. При использовании метода "Случайный лес" получилось добиться более высокого показателя точности **(0.873)**, поэтому в целом предпочтительнее использовать его.